

Lecture 4

OpenMP: Contents

- OpenMP's constructs fall into 5 categories:
 - ◆ Parallel Regions
 - ◆ Worksharing
 - ◆ Data Environment
 - ◆ Synchronization
 - ◆ Runtime functions/environment variables



OpenMP: Work-Sharing Constructs

- The “for” Work-Sharing construct splits up loop iterations among the threads in a team

```
#pragma omp parallel
#pragma omp for
    for (I=0;I<N;I++){
        NEAT_STUFF(I);
    }
```

By default, there is a barrier at the end of the “omp for”. Use the “nowait” clause to turn off the barrier.

Work Sharing Constructs

A motivating example

Sequential code

```
for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

OpenMP parallel region

```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    for(i=istart;i<iend;i++) { a[i] = a[i] + b[i];}
}
```

OpenMP parallel region and a work-sharing for-construct

```
#pragma omp parallel
#pragma omp for schedule(static)
    for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

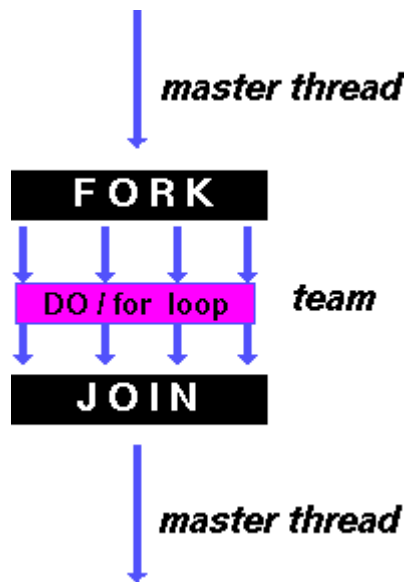


OpenMP Work-sharing constructs:

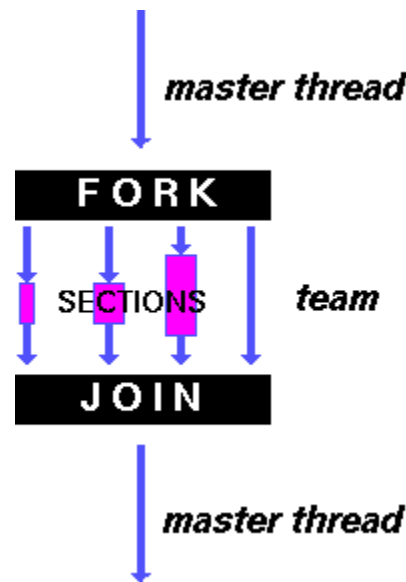
- A work-sharing construct divides the execution of the enclosed code region among the members of the team that encounter it.
- Work-sharing constructs do not launch new threads
- There is no implied barrier upon entry to a work-sharing construct, however there is an implied barrier at the end of a work sharing construct.

Types of Work-Sharing Constructs:

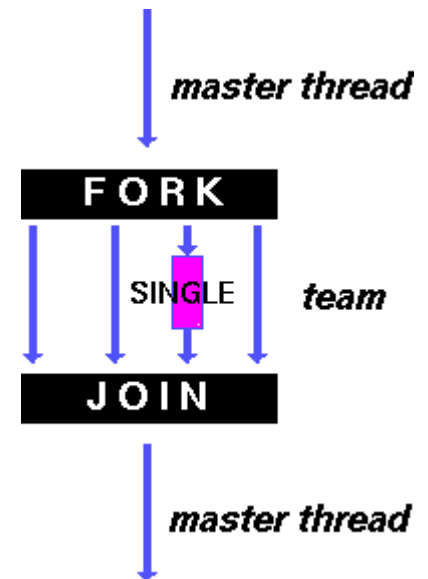
DO / for - shares iterations of a loop across the team.
Represents a type of "data parallelism".



SECTIONS - breaks work into separate, discrete sections.
Each section is executed by a thread. Can be used to implement a type of "functional parallelism".



SINGLE - serializes a section of code



OpenMP Work-sharing constructs: Notes

- A work-sharing construct must be enclosed dynamically within a parallel region in order for the directive to execute in parallel.
- Work-sharing constructs must be encountered by all members of a team or none at all
- The sequence of work-sharing regions and barrier regions encountered must be the same for every thread in a team.

Work-sharing constructs: Loop construct

- The DO / for directive specifies that the iterations of the loop immediately following it must be executed in parallel by the team. This assumes a parallel region has already been initiated, otherwise it executes in serial on a single processor.

```
#pragma omp parallel
```

```
#pragma omp for
```

```
  for (I=0;I<N;I++){  
    NEAT_STUFF(I);  
  }
```

```
!$omp parallel
```

```
!$omp do
```

```
  do I=1,N
```

```
    call NEAT_STUFF(I)
```

```
  end do
```

```
!$omp end do
```

```
!$omp end parallel
```


Simple examples: serial do-loop code

```
program loop
  implicit none
  integer, parameter :: N = 60000000
  integer :: i
  real :: x(N)

  do i = 1, N
    x(i) = 1./real(i)
  end do

end program
```

</home/syam/ces745/openmp/Fortran/do-loop/simple-do/loop-seq.f90>

parallel do-loop

One possible parallel version of the preceding code (distribute the loop to different threads by hard coding).

```
program loop
  use omp_lib
  implicit none
  integer, parameter :: N = 60000000
  integer :: i
  integer :: nprocs, myid, nb, istart, iend
  real :: x(N)
```

```
!$omp parallel private(myid,istart,iend)
```

```
  nprocs = omp_get_num_threads()
```

```
  myid = omp_get_thread_num()
```

```
  nb = N/nprocs
```

```
  istart = myid*nb + 1
```

```
  if (myid /= nprocs-1) then
```

```
    iend = (myid + 1)*nb
```

```
  else
```

```
    iend = N
```

```
  end if
```

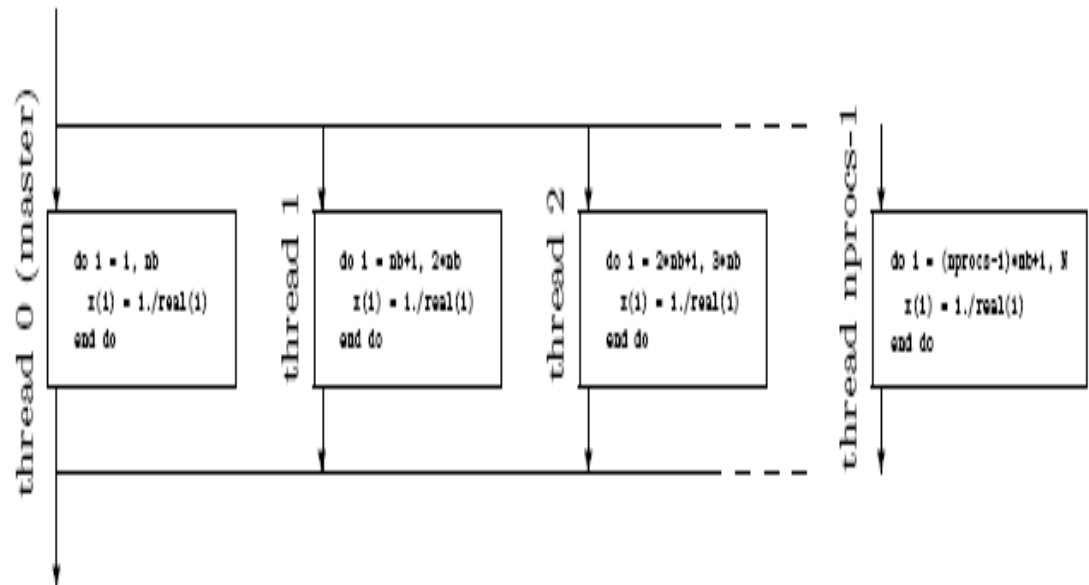
```
  do i = istart, iend
```

```
    x(i) = 1./real(i)
```

```
  end do
```

```
!$omp end parallel
```

```
end program
```



Do directive

Instead of hard-coding,
we can simply use
OpenMP loop directive
to achieve the same
goal.

```
program loop
  implicit none
  integer, parameter :: N = 60000000
  integer :: i
  real :: x(N)

  !$omp parallel
    !$omp do
      do i = 1, N
        x(i) = 1./real(i)
      end do
    !$omp end do
  !$omp end parallel

end program
```

Parallel do: Combined Directives

```
program loop
  implicit none
  integer, parameter :: N = 60000000
  integer :: i
  real :: x(N)

  !$omp parallel do
    do i = 1, N
      x(i) = 1./real(i)
    end do
  !$omp end parallel do

end program
```

DO/for Format

Fortran

```
!$OMP DO [clause ...]  
    SCHEDULE (type [,chunk])  
    ORDERED  
    PRIVATE (list)  
    FIRSTPRIVATE (list)  
    LASTPRIVATE (list)  
    SHARED (list)  
    REDUCTION (operator | intrinsic : list)  
do_loop  
!$OMP END DO [ NOWAIT ]
```

C/C++

```
#pragma omp for [clause ...] newline  
    schedule (type [,chunk])  
    ordered  
    private (list)  
    firstprivate (list)  
    lastprivate (list)  
    shared (list)  
    reduction (operator: list)  
    nowait  
for_loop
```

OpenMP For construct:

The schedule clause

- The schedule clause effects how loop iterations are mapped onto threads
 - ◆ `schedule(static [,chunk])`
 - Deal-out blocks of iterations of size “chunk” to each thread.
 - ◆ `schedule(dynamic[,chunk])`
 - Each thread grabs “chunk” iterations off a queue until all iterations have been handled.
 - ◆ `schedule(guided[,chunk])`
 - Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds.
 - ◆ `schedule(runtime)`
 - Schedule and chunk size taken from the `OMP_SCHEDULE` environment variable.

schedule(static)

- Iterations are divided evenly among threads

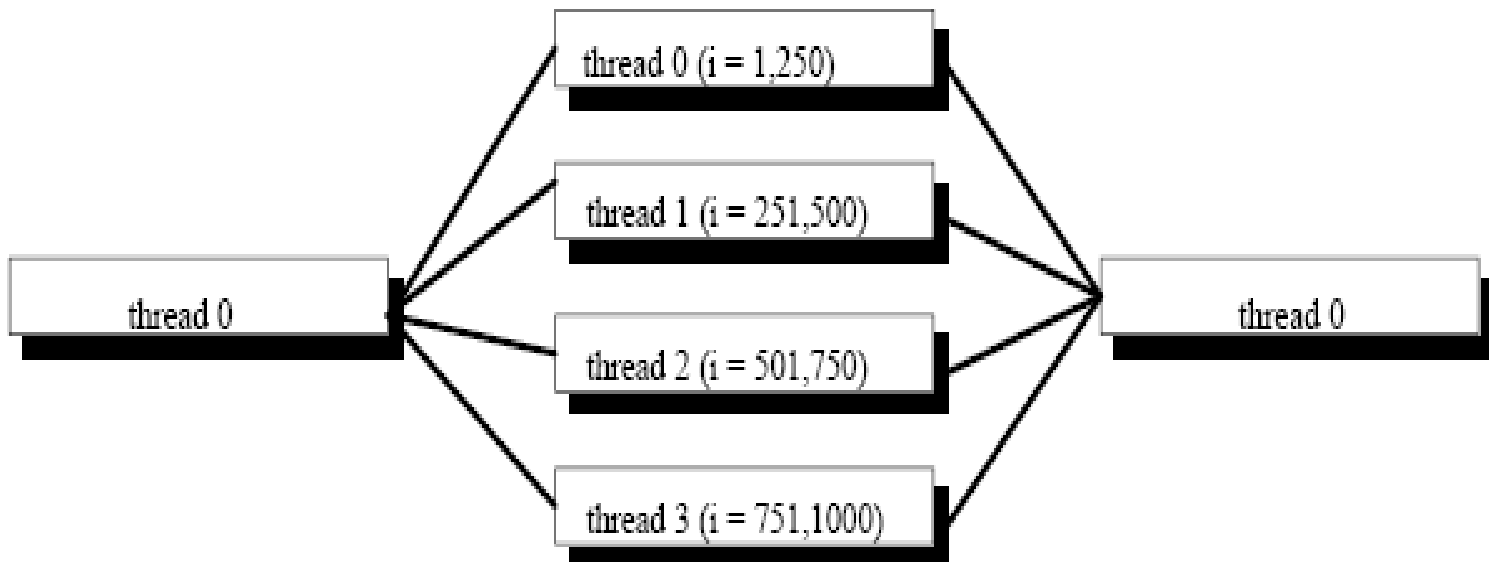
```
c$omp do shared(x) private(i)
```

```
c$omp& schedule(static)
```

```
do i = 1, 1000
```

```
x(i)=a
```

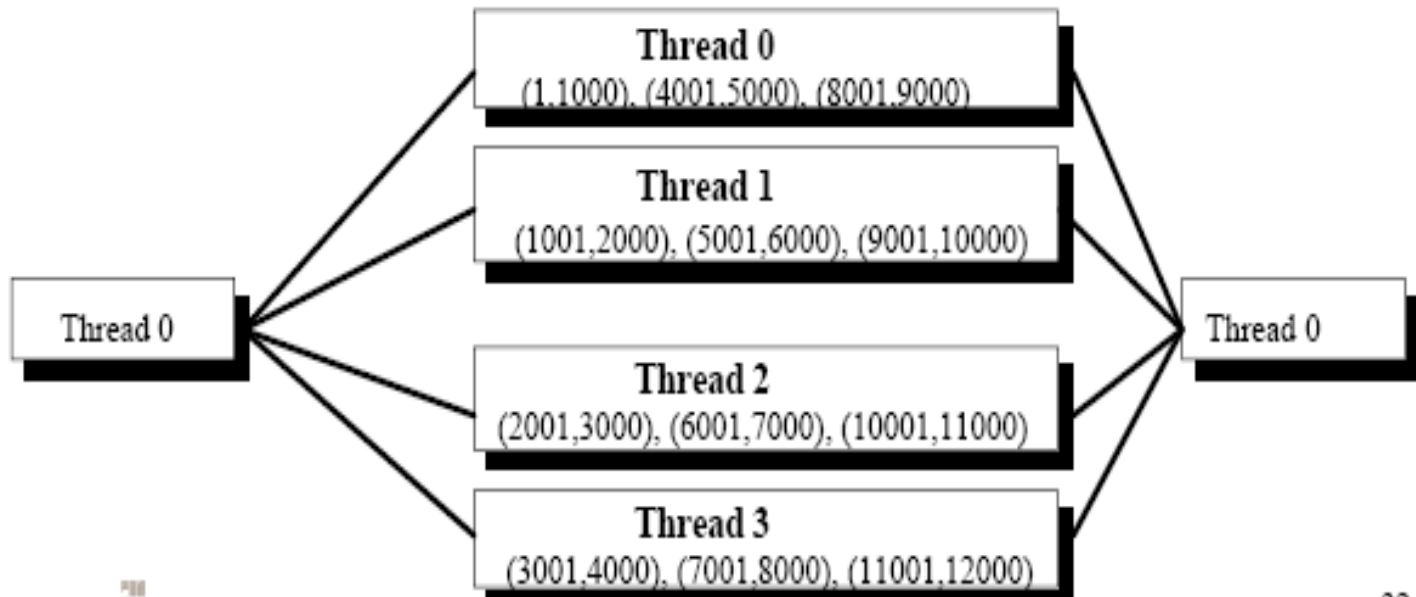
```
enddo
```



schedule(static,chunk)

- Divides the work load in to chunk sized parcels
- If there are N threads, each thread does every Nth chunk of work

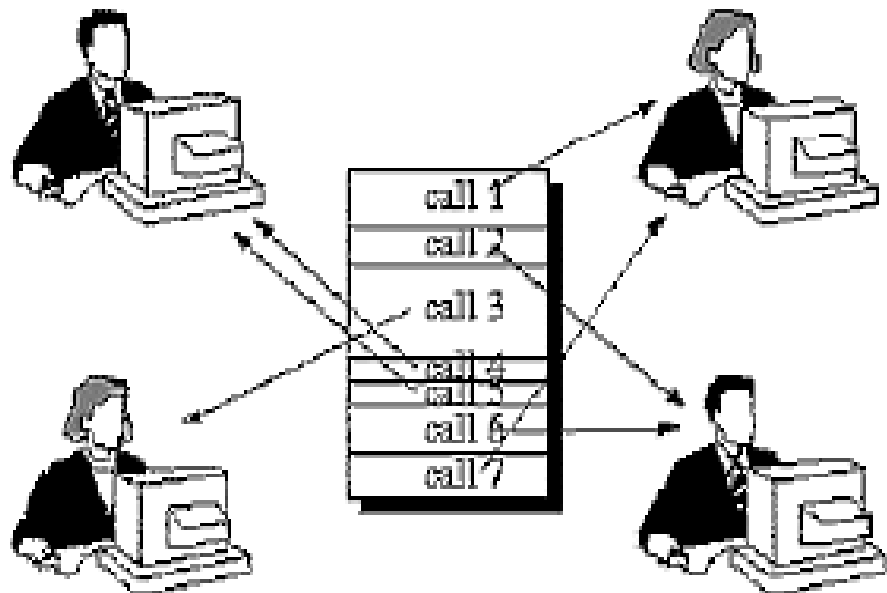
```
c$omp do shared(x)private(i)
c$omp& schedule(static,1000)
  do i = 1, 12000
    ... work ...
  enddo
```



schedule(dynamic,chunk)

- Divides the workload into chunk sized parcels.
- As a thread finishes one chunk, it grabs the next available chunk.
- Default value for chunk is 1.
- More overhead, but potentially better load balancing.

```
c$omp do shared(x) private(i)
c$omp& schedule(dynamic,1000)
  do i = 1, 10000
    ... work ...
  end do
```



schedule(guided,chunk)

- Like dynamic scheduling, but the chunk size varies dynamically.
- Chunk sizes depend on the number of unassigned iterations.
- The chunk size decreases toward the specified value of chunk.
- Achieves good load balancing with relatively low overhead.
- Insures that no single thread will be stuck with a large number of leftovers while the others take a coffee break.

```
c$omp do shared(x) private(i)
c$omp& schedule(guided,55)
do i = 1, 12000
... work ...
end do
```

schedule(runtime)

- Scheduling method is determined at runtime.
- Depends on the value of environment variable **OMP_SCHEDULE**
- This environment variable is checked at runtime, and the method is set accordingly.
- Scheduling method is static by default.
- Chunk size set as (optional) second argument of string expression.
- Useful for experimenting with different scheduling methods without recompiling.

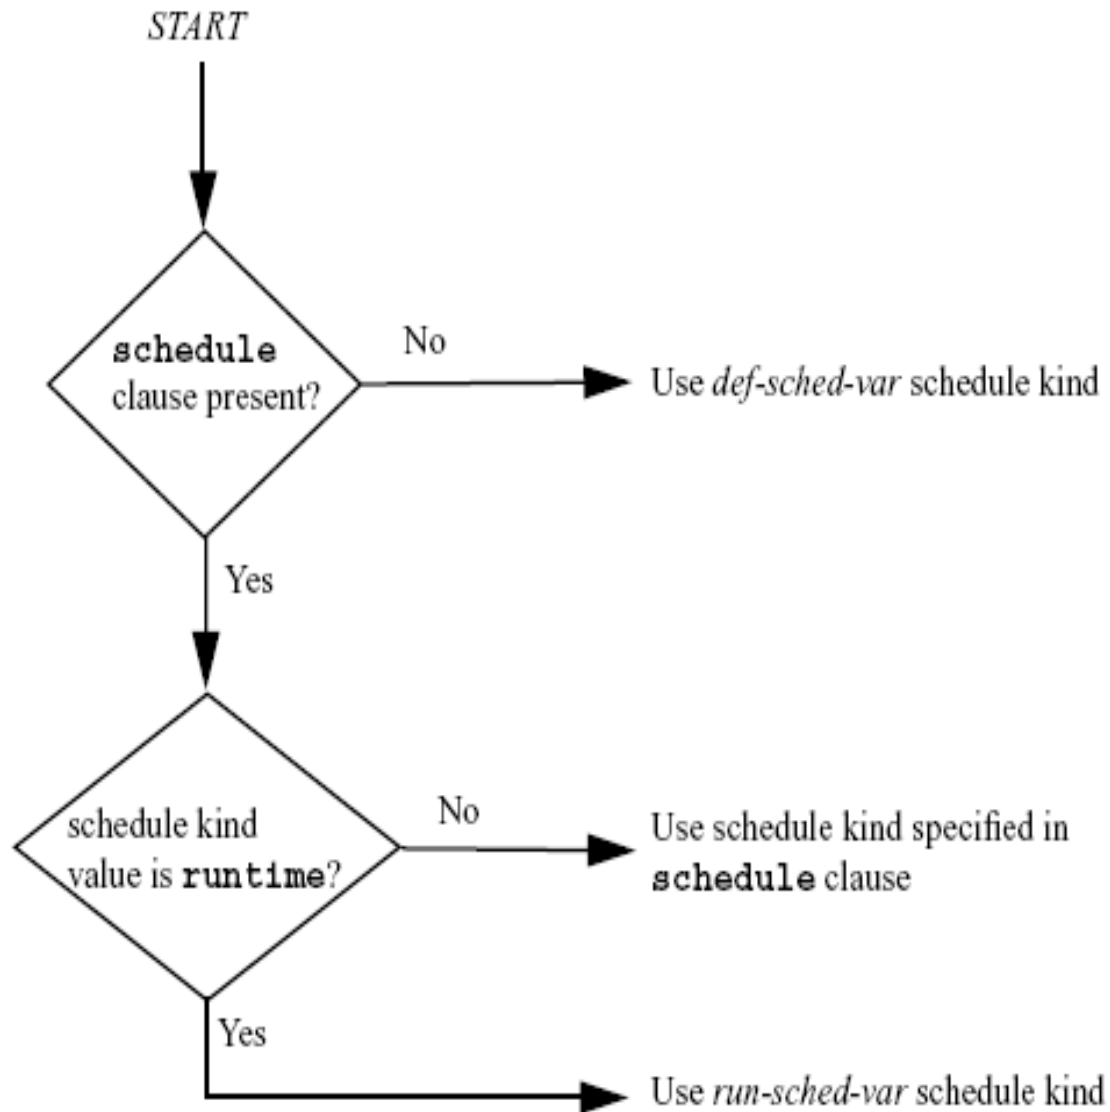
```
origin% export OMP_SCHEDULE=static,1000  
origin% export OMP_SCHEDULE=dynamic
```



DO/for construct: Notes

- The DO loop can not be a DO WHILE loop, or a loop without loop control. Also, the loop iteration variable must be an integer and the loop control parameters must be the same for all threads.
- Program correctness must not depend upon which thread executes a particular iteration.
- It is illegal to branch out of a loop associated with a DO/for directive.
- The chunk size must be specified as a loop invariant integer expression, as there is no synchronization during its evaluation by different threads.
- ORDERED and SCHEDULE clauses may appear once each.





Determining the schedule for a work-sharing loop.

Example: Simple vector-add program

- Three Arrays: A, B, C
- Arrays A, B, C, and variable N will be shared by all threads.
- Variable I will be private to each thread; each thread will have its own unique copy.
- The iterations of the loop will be distributed dynamically in CHUNK sized pieces.
- Threads will not synchronize upon completing their individual pieces of work (NOWAIT).

Fortran - DO Directive Example

```
PROGRAM VEC_ADD_DO
```

```
INTEGER N, CHUNKSIZE, CHUNK, I
```

```
PARAMETER (N=1000)
```

```
PARAMETER (CHUNKSIZE=100)
```

```
REAL A(N), B(N), C(N)
```

```
!      Some initializations
```

```
DO I = 1, N
```

```
    A(I) = I * 1.0
```

```
    B(I) = A(I)
```

```
ENDDO
```

```
CHUNK = CHUNKSIZE
```

```
!$OMP PARALLEL SHARED(A,B,C,CHUNK) PRIVATE(I)
```

```
!$OMP DO SCHEDULE(DYNAMIC,CHUNK)
```

```
DO I = 1, N
```

```
    C(I) = A(I) + B(I)
```

```
ENDDO
```

```
!$OMP END DO NOWAIT
```

```
!$OMP END PARALLEL
```

```
END
```

C / C++ - for Directive Example

```
#include <omp.h>
#define CHUNKSIZE 100
#define N 1000

main () {
    int i, chunk;
    float a[N], b[N], c[N];

    /* Some initializations */
    for (i=0; i < N; i++)
        a[i] = b[i] = i * 1.0;
    chunk = CHUNKSIZE;

    #pragma omp parallel shared(a,b,c,chunk) private(i)
    {
        #pragma omp for schedule(dynamic,chunk) nowait
        for (i=0; i < N; i++)
            c[i] = a[i] + b[i];
    } /* end of parallel section */

}
```


More on Schedule clause, demo

</home/syam/ces745/openmp/Fortran/do-loop/schedule>

do-schedule-static.f

do-schedule-dynamic.f

do-schedule-guide.f

do-schedule-runtime.f

Work-Sharing Constructs: SECTIONS Directive

Purpose:

The SECTIONS directive is a non-iterative work-sharing construct. It specifies that the enclosed section(s) of code are to be divided among the threads in the team.

Independent SECTION directives are nested within a SECTIONS directive. Each SECTION is executed once by a thread in the team. Different sections may be executed by different threads. It is possible for a thread to execute more than one section if it is quick enough and the implementation permits such.

OpenMP: Work-Sharing Constructs

- The Sections work-sharing construct gives a different structured block to each thread.

```
#pragma omp parallel
#pragma omp sections
{
    X_calculation();
#pragma omp section
    y_calculation();
#pragma omp section
    z_calculation();
}
```

By default, there is a barrier at the end of the “omp sections”. Use the “nowait” clause to turn off the barrier.

Format:

Fortran

```
!$OMP SECTIONS [clause ...]  
    PRIVATE (list)  
    FIRSTPRIVATE (list)  
    LASTPRIVATE (list)  
    REDUCTION (operator | intrinsic : list)  
!$OMP SECTION  
    structured_block  
!$OMP SECTION  
    structured_block  
!$OMP END SECTIONS [ NOWAIT ]
```

C/C++

```
#pragma omp sections [clause ...] newline  
    private (list)  
    firstprivate (list)  
    lastprivate (list)  
    reduction (operator: list)  
    nowait  
{  
#pragma omp section newline  
    structured_block  
#pragma omp section newline  
    structured_block  
}
```

Clauses:

1. There is an implied barrier at the end of a `SECTIONS` directive, unless the `NOWAIT/nowait` clause is used.
2. Clauses are described in detail later, in the Data Scope Attribute section.

Restrictions:

1. It is illegal to branch into or out of section blocks.
2. `SECTION` directives must occur within the lexical extent of an enclosing `SECTIONS` directive

Examples: 3-loops

Serial code with three independent tasks, namely, three do loops, each operating on a different array using different loop counters and temporary scalar variables.

```
program compute
  implicit none
  integer, parameter :: NX = 10000000
  integer, parameter :: NY = 20000000
  integer, parameter :: NZ = 30000000
  real :: x(NX)
  real :: y(NY)
  real :: z(NZ)
  integer :: i, j, k
  real :: ri, rj, rk
  write(*,*) "start"
  do i = 1, NX
    ri = real(i)
    x(i) = atan(ri)/ri
  end do
  do j = 1, NY
    rj = real(j)
    y(j) = cos(rj)/rj
  end do
  do k = 1, NZ
    rk = real(k)
    z(k) = log10(rk)/rk
  end do
  write(*,*) "end"
end program
```

Examples: 3-loops

One possible parallel version of the preceding code (distribute the loop to different threads by hard coding).

program compute

.....

write(*,*) "start"

!\$omp parallel

select case (omp_get_thread_num())

case (0)

do i = 1, NX

ri = real(i)

x(i) = atan(ri)/ri

end do

case (1)

do j = 1, NY

rj = real(j)

y(j) = cos(rj)/rj

end do

case (2)

do k = 1, NZ

rk = real(k)

z(k) = log10(rk)/rk

end do

end select

!\$omp end parallel

write(*,*) "end"

end program

Examples: 3-loops

Instead of hard-coding,
we can use OpenMP
task sharing directives
(section) to achieve the
same goal.

program compute

.....

write(*,*) "start"

!\$omp parallel

!\$omp sections

!\$omp section

do i = 1, NX

ri = real(i)

x(i) = atan(ri)/ri

end do

!\$omp section

do j = 1, NY

rj = real(j)

y(j) = cos(rj)/rj

end do

!\$omp section

do k = 1, NZ

rk = real(k)

z(k) = log10(rk)/rk

end do

!\$omp end sections

!\$omp end parallel

write(*,*) "end"

end program

More examples:

Combined Directives

```
program compute
```

```
.....
```

```
write(*,*) "start"
```

```
!$omp parallel sections
```

```
!$omp section
```

```
do i = 1, NX
```

```
ri = real(i)
```

```
x(i) = atan(ri)/ri
```

```
end do
```

```
!$omp section
```

```
do j = 1, NY
```

```
rj = real(j)
```

```
y(j) = cos(rj)/rj
```

```
end do
```

```
!$omp section
```

```
do k = 1, NZ
```

```
rk = real(k)
```

```
z(k) = log10(rk)/rk
```

```
end do
```

```
!$omp end parallel sections
```

```
write(*,*) "end"
```

```
end program
```

Example: Vector-add

Fortran: vector-add

```
PROGRAM VEC_ADD_SECTIONS
INTEGER N, I
PARAMETER (N=1000)
REAL A(N), B(N), C(N)
! Some initializations
DO I = 1, N
    A(I) = I * 1.0
    B(I) = A(I)
ENDDO
```

```
!$OMP PARALLEL SHARED(A,B,C)
    PRIVATE(I)
!$OMP SECTIONS
!$OMP SECTION
    DO I = 1, N/2
        C(I) = A(I) + B(I)
    ENDDO
!$OMP SECTION
    DO I = 1+N/2, N
        C(I) = A(I) + B(I)
    ENDDO
!$OMP END SECTIONS NOWAIT
!$OMP END PARALLEL
END
```

- The first $n/2$ iterations of the DO loop will be distributed to the first thread, and the rest will be distributed to the second thread
- When each thread finishes its block of iterations, it proceeds with whatever code comes next (NOWAIT)

C/C++: vector-add

```
#include <omp.h>
#define N 1000
```

```
main () {
int i; float a[N], b[N], c[N];
/* Some initializations */
for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;
```

```
#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp sections nowait
    {
        #pragma omp section
        for (i=0; i < N/2; i++)
            c[i] = a[i] + b[i];
        #pragma omp section
        for (i=N/2; i < N; i++)
            c[i] = a[i] + b[i];
    } /* end of sections */
} /* end of parallel section */
}
```

Questions:

What happens if the number of threads and the number of SECTIONS are different? More threads than SECTIONS? Fewer threads than SECTIONS?

Answer: If there are more threads than sections, some threads will not execute a section and some will. If there are more sections than threads, the implementation defines how the extra sections are executed.

Which thread executes which SECTION?

Answer: It is up to the implementation to decide which threads will execute a section and which threads will not, and it can vary from execution to execution

Work-Sharing Constructs: **SINGLE Directive**

Purpose:

The SINGLE directive specifies that the enclosed code is to be executed by only one thread in the team.

May be useful when dealing with sections of code that are not thread safe (such as I/O)

OpenMP Work Sharing Constructs - single

- Ensures that a code block is executed by only one thread in a parallel region.
- The thread that reaches the single directive first is the one that executes the single block.
- Equivalent to a sections directive with a single section - but a more descriptive syntax.
- All threads in the parallel region must encounter the single directive.
- Unless nowait is specified, all uninvolved threads wait at the end of the single block

```
c$omp parallel private(i) shared(a)
c$omp do
    do i = 1, n
        ...work on a(i) ...
    enddo
c$omp single
    ... process result of do ...
c$omp end single
c$omp do
    do i = 1, n
        ... more work ...
    enddo
c$omp end parallel
```

OpenMP Work Sharing Constructs - single

- Fortran syntax:

```
c$omp single [clause [clause...]]  
    structured block  
c$omp end single [nowait]
```

where clause is one of

- private(*list*)
- firstprivate(*list*)

OpenMP Work Sharing Constructs - single

- C syntax:

```
#pragma omp single [clause [clause...]]  
    structured block
```

where clause is one of

- private(*list*)
- firstprivate(*list*)
- nowait

OpenMP Work Sharing Constructs - single

Clauses:

- Threads in the team that do not execute the SINGLE directive, wait at the end of the enclosed code block, unless a NOWAIT/nowait clause is specified.

Restrictions:

- It is illegal to branch into or out of a SINGLE block.

Examples

```
PROGRAM single_1
```

```
write(*,*) 'start'
```

```
!$OMP PARALLEL DEFAULT(NONE), private(i)
```

```
!$OMP DO
```

```
do i=1,5
```

```
write(*,*) i
```

```
enddo
```

```
!$OMP END DO
```

```
!$OMP SINGLE
```

```
write(*,*) 'begin single directive'
```

```
do i=1,5
```

```
write(*,*) 'hello',i
```

```
enddo
```

```
!$OMP END SINGLE
```

```
!$OMP END PARALLEL
```

```
write(*,*) 'end'
```

```
END
```

```
[syam@wha780 single]$ ./single-1
```

```
start
```

```
1
```

```
4
```

```
5
```

```
2
```

```
3
```

```
begin single directive
```

```
hello 1
```

```
hello 2
```

```
hello 3
```

```
hello 4
```

```
hello 5
```

```
end
```

```
[syam@wha780 single]$
```

```
PROGRAM single_3
INTEGER NTHREADS, TID, TID2,
OMP_GET_NUM_THREADS, OMP_GET_THREAD_NUM
```

```
write(*,*) "Start"
```

```
!$OMP PARALLEL PRIVATE(TID, i)
```

```
!$OMP DO
```

```
do i=1,8
```

```
TID = OMP_GET_THREAD_NUM()
```

```
write(*,*) "thread: ", TID, 'i = ', i
```

```
enddo
```

```
!$OMP END DO
```

```
!$OMP SINGLE
```

```
write(*,*) "SINGLE - begin"
```

```
do i=1,8
```

```
TID2 = OMP_GET_THREAD_NUM()
```

```
PRINT *, 'This is from thread = ', TID2
```

```
write(*,*) 'hello',i
```

```
enddo
```

```
!$OMP END SINGLE
```

```
!$OMP END PARALLEL
```

```
write(*,*) "End "
```

```
END
```

```
[syam@wha780 single]$
```

```
./single-3
```

```
Start
```

```
thread: 0 i = 1
```

```
thread: 1 i = 5
```

```
thread: 1 i = 6
```

```
thread: 1 i = 7
```

```
thread: 1 i = 8
```

```
thread: 0 i = 2
```

```
thread: 0 i = 3
```

```
thread: 0 i = 4
```

```
SINGLE - begin
```

```
This is from thread = 0
```

```
hello 1
```

```
This is from thread = 0
```

```
hello 2
```

```
This is from thread = 0
```

```
hello 3
```

```
This is from thread = 0
```

```
hello 4
```

```
This is from thread = 0
```

```
hello 5
```

```
This is from thread = 0
```

```
hello 6
```

```
This is from thread = 0
```

```
hello 7
```

```
This is from thread = 0
```

```
hello 8
```

```
End
```

OpenMP: Combined Parallel Work-Sharing Constructs

- A short hand notation that combines the Parallel and work-sharing construct.

```
#pragma omp parallel for  
    for (l=0;l<N;l++){  
        NEAT_STUFF(l);  
    }
```

- There's also a “parallel sections” construct.

OpenMP:

More details: Scope of OpenMP constructs

OpenMP constructs can span multiple source files.

